

MPI Send and Receive Bcast Operations

Introduction to MPI Communication

MPI (Message Passing Interface) is a standardized API for communicating between processes in parallel computing. The two fundamental operations are:

1. **MPI_Send** - sends a message to another process
2. **MPI_Recv** - receives a message from another process

MPI_Send - Sending Data

```
int MPI_Send(  
    const void *buf,           // Address of the data to send  
    int count,                 // Number of elements to send  
    MPI_Datatype datatype,     // Type of each element  
    int dest,                  // Rank of destination process  
    int tag,                   // Message tag (integer)  
    MPI_Comm comm              // Communicator (usually MPI_COMM_WORLD)  
);
```

Example of MPI_Send

```
int data = 42;  
MPI_Send(&data, 1, MPI_INT, 1, 0, MPI_COMM_WORLD);
```

This sends the integer value 42 to process with rank 1, with tag 0.

MPI_Recv - Receiving Data

```
int MPI_Recv(  
    void *buf,                 // Address to store received data  
    int count,                 // Maximum number of elements to receive  
    MPI_Datatype datatype,     // Type of each element  
    int source,                // Rank of sending process (or MPI_ANY_SOURCE)  
    int tag,                   // Message tag (or MPI_ANY_TAG)  
    MPI_Comm comm,             // Communicator  
    MPI_Status *status         // Status object with info about the message  
);
```

Example of MPI_Recv

```
int received_data;
MPI_Status status;
MPI_Recv(&received_data, 1, MPI_INT, 0, 0, MPI_COMM_WORLD, &status);
```

This receives an integer from process 0 with tag 0 and stores it in `received_data`.

Complete Example: Sending and Receiving Between Two Processes

```
#include <mpi.h>
#include <stdio.h>

int main(int argc, char** argv) {
    MPI_Init(&argc, &argv);

    int world_rank;
    MPI_Comm_rank(MPI_COMM_WORLD, &world_rank);

    int world_size;
    MPI_Comm_size(MPI_COMM_WORLD, &world_size);

    if (world_size < 2) {
        printf("This program requires at least 2 processes\n");
        MPI_Finalize();
        return 1;
    }

    if (world_rank == 0) {
        // Process 0 sends a message to process 1
        int number = 12345;
        MPI_Send(&number, 1, MPI_INT, 1, 0, MPI_COMM_WORLD);
        printf("Process 0 sent number %d to process 1\n", number);
    } else if (world_rank == 1) {
        // Process 1 receives the message from process 0
        int number;
        MPI_Status status;
        MPI_Recv(&number, 1, MPI_INT, 0, 0, MPI_COMM_WORLD, &status);
        printf("Process 1 received number %d from process 0\n", number);
    }
}
```

```
MPI_Finalize();  
return 0;  
}
```

Key Points

1. **Blocking Nature:** Both `MPI_Send` and `MPI_Recv` are blocking calls - they don't return until the operation is complete.
2. **Matching Rules:**
 - The sender's `dest` must match the receiver's `source` (or use `MPI_ANY_SOURCE`)
 - The tags must match (or use `MPI_ANY_TAG`)
 - The datatypes must match
 - The communicator must match
3. **Status Object:** Contains information about the received message:
 - `status.MPI_SOURCE` - who sent the message
 - `status.MPI_TAG` - what tag was used
 - Can query message length with `MPI_Get_count`
4. **Common Pitfalls:**
 - Deadlocks (both processes waiting to receive before sending)
 - Mismatched tags/datatypes
 - Buffer overflow (receiving more data than buffer can hold)

Beyond the basic `MPI_Send` and `MPI_Recv`, MPI offers several other communication methods that are useful for different scenarios in parallel programming. Here's a comprehensive overview:

1. Non-blocking Communication

`MPI_Isend` (Immediate Send)

```
int MPI_Isend(
    const void *buf,
    int count,
    MPI_Datatype datatype,
    int dest,
    int tag,
    MPI_Comm comm,
    MPI_Request *request
);
```

- Returns immediately without waiting for the message to be sent
- The `request` object is used to check completion status

MPI_Irecv (Immediate Receive)

```
int MPI_Irecv(
    void *buf,
    int count,
    MPI_Datatype datatype,
    int source,
    int tag,
    MPI_Comm comm,
    MPI_Request *request
);
```

- Posts a receive operation that will complete later
- Allows overlap of computation and communication

Completion Checking

```
int MPI_Test(MPI_Request *request, int *flag, MPI_Status *status);
int MPI_Wait(MPI_Request *request, MPI_Status *status);
```

- `MPI_Test`: Checks if operation is complete (non-blocking)
- `MPI_Wait`: Waits until operation is complete (blocking)

Complete Example

```
#include <mpi.h>
#include <stdio.h>

int main(int argc, char** argv) {
    MPI_Init(&argc, &argv);

    int rank, size;
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &size);

    if (size < 2) {
        printf("Need at least 2 processes\n");
        MPI_Finalize();
        return 1;
    }

    int data_to_send = rank * 100;
    int received_data;
    MPI_Request send_request, recv_request;
    MPI_Status status;

    if (rank == 0) {
        // Process 0 sends to 1 and receives from 1
        MPI_Isend(&data_to_send, 1, MPI_INT, 1, 0, MPI_COMM_WORLD, &send_request);
        MPI_Irecv(&received_data, 1, MPI_INT, 1, 0, MPI_COMM_WORLD, &recv_request);
;

        // Do some computation while communication happens
        printf("Process 0 is computing while waiting...\n");
        for (int i = 0; i < 1000000; i++) {} // Simulate work

        // Wait for both operations to complete
        MPI_Wait(&send_request, MPI_STATUS_IGNORE);
        MPI_Wait(&recv_request, &status);

        printf("Process 0 received %d from process 1\n", received_data);
    }
    else if (rank == 1) {
        // Process 1 sends to 0 and receives from 0
        MPI_Isend(&data_to_send, 1, MPI_INT, 0, 0, MPI_COMM_WORLD, &send_request);
        MPI_Irecv(&received_data, 1, MPI_INT, 0, 0, MPI_COMM_WORLD, &recv_request);
;

        // Do some computation while communication happens
        printf("Process 1 is computing while waiting...\n");
        for (int i = 0; i < 1000000; i++) {} // Simulate work

        // Wait for both operations to complete
```

```

    MPI_Wait(&send_request, MPI_STATUS_IGNORE);
    MPI_Wait(&recv_request, &status);

    printf("Process 1 received %d from process 0\n", received_data);
}

MPI_Finalize();
return 0;
}

```

Important Considerations

1. Buffer Management:

- Don't modify send buffers until send completes
- Don't access receive buffers until receive completes
- Use separate buffers for overlapping communications

2. Request Objects:

- Each non-blocking operation needs its own request object
- Request objects must be matched with completion calls
- Never reuse a request object before its operation completes

3. Performance Benefits:

- Greatest when there's significant computation to overlap
- Less beneficial for very short messages

4. Error Handling:


- Always check return codes in production code
- Use `MPI_Status` to check for message details

3. Collective Communication

Understanding and Using MPI_Bcast

What is MPI_Bcast?

MPI_Bcast (Broadcast) is a **collective communication operation** that allows one process (the root) to send the same data to all other processes in a communicator. It's one of the most commonly used MPI functions for distributing data to all processes.



Broadcast (MPI_Bcast)

```
int MPI_Bcast(
    void *buffer,
    int count,
    MPI_Datatype datatype,
    int root,
    MPI_Comm comm
);
```

- One process (root) sends the same data to all processes

Why Use MPI_Bcast?

1. **Efficiency:**

- More optimized than sending individually to each process
- Uses tree-based algorithms (logarithmic complexity) rather than linear sends

2. **Synchronization:**

- All processes reach this point together (implicit barrier)
- Ensures all processes have the data before proceeding

3. **Simplicity:**

- Single call replaces multiple send/receive operations

- Cleaner, more readable code

How MPI_Bcast Works

1. The **root process** provides the data in its buffer
2. All other processes provide an empty buffer of sufficient size
3. After completion, all processes have identical data in their buffers

```
#include <mpi.h>
#include <stdio.h>
#include <stdlib.h>

int main(int argc, char** argv) {
    MPI_Init(&argc, &argv);

    int rank, size;
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &size);

    const int ROWS = 3, COLS = 4;
    double matrix[ROWS][COLS];

    // Root process initializes the matrix
    if (rank == 0) {
        printf("Root initializing matrix:\n");
        for (int i = 0; i < ROWS; i++) {
            for (int j = 0; j < COLS; j++) {
                matrix[i][j] = i * 10 + j;
                printf("%.1f ", matrix[i][j]);
            }
            printf("\n");
        }
    }

    // Broadcast the entire matrix to all processes
    MPI_Bcast(matrix, ROWS*COLS, MPI_DOUBLE, 0, MPI_COMM_WORLD);

    MPI_Barrier(MPI_COMM_WORLD);
    for (int p = 0; p < size; p++) {
        if (rank == p) {
            printf("Process %d received:\n", rank);
            for (int i = 0; i < ROWS; i++) {
                for (int j = 0; j < COLS; j++) {
                    printf("%.1f ", matrix[i][j]);
                }
            }
        }
    }
}
```



```
        printf("\n");
    }
    fflush(stdout);
}
MPI_Barrier(MPI_COMM_WORLD);
}
MPI_Finalize();
return 0;
}
```