


Week11

MPI_Scatter and MPI_Gather

These are two fundamental **collective communication** operations in MPI used for **distributing** and **collecting** data among processes.

1. MPI_Scatter - Distribute Data from Root to All Processes



Data in an array on root node:

A(0) A(1) A(2) A(N-1)

Goes to processors:

P₀ P₁ P₂ . . . P_{n-1}

Prototype

```
int MPI_Scatter(  
    const void* sendbuf,    // Buffer at root containing all data (input, significant only at root)  
    int sendcount,          // Number of elements sent to each process  
    MPI_Datatype sendtype,  // Data type of sent elements (e.g., MPI_INT)  
    void* recvbuf,         // Buffer where received data will be stored (output)  
    int recvcount,         // Number of elements to receive (should match sendcount)  
    MPI_Datatype recvtype,  // Data type of received elements  
    int root,              // Rank of the root process (the one that scatters)  
    MPI_Comm comm          // Communicator (usually MPI_COMM_WORLD)  
);
```

Why Use MPI_Scatter?

- **Distributes a large dataset** from the **root process** to **all other processes**.
- **Balanced workload**: Each process gets an equal chunk of data.
- **Efficient alternative** to manually sending data with MPI_Send/MPI_Recv.

Example


```
int data[12] = {1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12}; // Only root has this  
int local_data[3]; // Each process gets 3 elements  
  
MPI_Scatter(  


```

```
data,          // Root sends this entire array
3,            // Each process gets 3 elements
MPI_INT,       // Data type
local_data,    // Each process stores its part here
3,            // Each process expects 3 elements
MPI_INT,
0,            // Root (rank 0) does the scattering
MPI_COMM_WORLD
);

// Now:
// - Process 0 gets {1, 2, 3}
// - Process 1 gets {4, 5, 6}
// - Process 2 gets {7, 8, 9}
// - Process 3 gets {10, 11, 12}
```

2. MPI_Gather - Collect Data from All Processes to Root





Prototype

```
int MPI_Gather(
    const void* sendbuf,    // Data to be sent by each process (input)
    int sendcount,         // Number of elements sent by each process
    MPI_Datatype sendtype, // Data type of sent elements
    void* recvbuf,         // Buffer at root to store gathered data (output, significant only at root)
    int recvcnt,           // Number of elements received per process (should match sendcount)
    MPI_Datatype recvtype, // Data type of received elements
    int root,              // Rank of the root process (collects data)
    MPI_Comm comm          // Communicator (usually MPI_COMM_WORLD)
);
```

Why Use MPI_Gather?

- **Combines results** from all processes into **a single array at the root**.
- **Useful after parallel computation** (e.g., each process computes a part, then gathers results).
- **More efficient** than manually collecting data with MPI_Send/MPI_Recv.

Example

```
int local_data[3] = {rank+1, rank+2, rank+3}; // Each process has its own data
int gathered_data[12]; // Only root will store the full result
```

```

MPI_Gather(
    local_data,      // Each process sends its 3 elements
    3,              // Each process sends 3 elements
    MPI_INT,
    gathered_data,   // Root collects all data here
    3,              // Root expects 3 elements from each process
    MPI_INT,
    0,              // Root (rank 0) collects data
    MPI_COMM_WORLD
);

// If rank == 0, gathered_data will be:
// {1, 2, 3, 2, 3, 4, 3, 4, 5, 4, 5, 6} (assuming 4 processes)

```

Key Differences Between MPI_Scatter and MPI_Gather

Feature	MPI_Scatter	MPI_Gather
Direction	Root → All	All → Root
Use Case	Distribute data for parallel work	Collect results after computation
Buffer Usage	sendbuf (root only), recvbuf (all)	sendbuf (all), recvbuf (root only)

When to Use Them?

- MPI_Scatter
 - Splitting a large dataset for parallel processing.
 - Example: Distributing rows of a matrix to different processes.
- MPI_Gather
 - Combining partial results after computation.

- Example: Collecting local sums into a global sum at the root.

Summary

- `MPI_Scatter` = Distribute data (one-to-many).
- `MPI_Gather` = Collect data (many-to-one).
- Both are **collective operations** (all processes must call them).
- **More efficient** than point-to-point communication (`MPI_Send`/`MPI_Recv`) for structured data distribution.

These operations are **essential for parallel algorithms** where work is divided among processes!



1. MPI_Scatter Example (Distribute Data)

Objective: Root process splits an array and sends chunks to all processes.

```
#include <stdio.h>
#include <mpi.h>

int main() {
    MPI_Init(NULL, NULL);

    int rank, size;
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &size);

    int data[12] = {10, 20, 30, 40, 50, 60, 70, 80, 90, 100, 110, 120}; // Only root has this
    int local_data[3]; // Each process will receive 3 elements

    // Scatter the data from root (rank 0)
```

```

MPI_Scatter(
    data, 3, MPI_INT,      // Send 3 ints from root's 'data'
    local_data, 3, MPI_INT, // Each process receives 3 ints into 'local_data'

    0, MPI_COMM_WORLD      // Root is rank 0
);

// Each process prints its received data
printf("Process %d received: %d, %d, %d\n",
    rank, local_data[0], local_data[1], local_data[2]);

MPI_Finalize();
return 0;
}

```

Output (for 4 processes):

```

Process 0 received: 10, 20, 30
Process 1 received: 40, 50, 60
Process 2 received: 70, 80, 90
Process 3 received: 100, 110, 120

```

2. MPI_Gather Example (Collect Data)

Objective: All processes send their data to the root for aggregation.

```

#include <stdio.h>
#include <mpi.h>

int main() {
    MPI_Init(NULL, NULL);

    int rank, size;
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &size);

    int local_data[2] = {rank * 10, rank * 10 + 5}; // Each process has 2 values
}

```

```

int gathered_data[8]; // Root will collect all data (4 processes × 2 elements)

// Gather data at root (rank 0)
MPI_Gather(
    local_data, 2, MPI_INT,    // Each process sends 2 ints
    gathered_data, 2, MPI_INT, // Root receives 2 ints per process
    0, MPI_COMM_WORLD         // Root is rank 0
);

// Root prints the gathered data
if (rank == 0) {
    printf("Gathered data at root: ");
    for (int i = 0; i < 8; i++) {
        printf("%d ", gathered_data[i]);
    }
    printf("\n");
}

MPI_Finalize();
return 0;
}

```

Output (for 4 processes):

Gathered data at root: 0 5 10 15 20 25 30 35

Key Notes

1. **MPI_Scatter:**
 - Root divides an array and sends chunks to all processes.
 - Non-root processes only need a receive buffer.
2. **MPI_Gather:**
 - All processes send data to the root.
 - Only root needs a receive buffer.

MPI_Reduce

is a collective communication operation in MPI (Message Passing Interface) that combines values from all processes in a communicator using a specified operation (like sum, max, min, etc.) and stores the result in a single target process (called the *root* process).

Prototype of MPI_Reduce

```
int MPI_Reduce(  
    const void* sendbuf,    // Address of the local data to be reduced (input)  
    void* recvbuf,         // Address where the reduced result will be stored  
                           // (output, significant only at root)  
    int count,             // Number of elements in the send buffer  
    MPI_Datatype datatype, // Data type of the elements (e.g., MPI_INT, MPI_FL  
OAT)  
    MPI_Op op,             // Reduction operation (e.g., MPI_SUM, MPI_MAX)  
    int root,              // Rank of the process that receives the result  
    MPI_Comm comm          // Communicator (usually MPI_COMM_WORLD)  
);
```

Why is MPI_Reduce Used?

1. Aggregates Data Efficiently

- Instead of manually gathering data to one process and then computing the result, `MPI_Reduce` performs the reduction in a single optimized step.
- Example: Summing values from all processes (`MPI_SUM`), finding the maximum (`MPI_MAX`), or computing logical AND (`MPI_LAND`).

2. Supports Various Reduction Operations

- Common operations:
 - `MPI_SUM` (Summation)
 - `MPI_PROD` (Product)
 - `MPI_MAX` (Maximum value)
 - `MPI_MIN` (Minimum value)
 - `MPI LAND` (Logical AND)
 - `MPI_BAND` (Bitwise AND)
 - And more...

3. Better Performance

- MPI implementations optimize `MPI_Reduce` for the underlying hardware (e.g., using tree-based algorithms for scalability).

4. Simplifies Parallel Computations

- Used in many numerical algorithms (e.g., dot products, matrix-vector multiplication, global statistics).

Example Breakdown

1. Each Process Computes a Local Sum

```
int a = rand() % 100;      // Random number 1
int b = rand() % 100;      // Random number 2
int local_sum = a + b;     // Local sum (different for each process)
```

2. Reduction to a Single Process (Root)

```
MPI_Reduce(
    &local_sum,      // Input: Each process sends its local_sum
    &global_sum,     // Output: Only root gets the final sum
    1,              // Only 1 value is being reduced
    MPI_INT,        // Data type is integer
```

```
MPI_SUM,          // Operation: Summation
0,               // Root process (rank 0 gets the result)
MPI_COMM_WORLD   // Use all processes
);
```

3. Root Process Prints the Final Result

```
if (rank == 0) {
    printf("Total sum from all processes: %d\n", global_sum);
}
```

Key Notes

- **MPI_Reduce vs MPI_Allreduce**
 - **MPI_Reduce** sends the result to **only one process (root)**.
 - **MPI_Allreduce** sends the result to **all processes** (useful when every process needs the reduced value).
- **Blocking Operation**
 - All processes must call **MPI_Reduce** (it synchronizes).
- **Efficiency**
 - Prefer **MPI_Reduce** over manual gather + compute for better performance.

Common Use Cases

1. **Parallel Summation** (e.g., computing total energy in simulations)
2. **Finding Global Min/Max** (e.g., optimization problems)
3. **Dot Products** (used in linear algebra)
4. **Statistical Aggregation** (e.g., computing mean/variance)

This makes `MPI_Reduce` a fundamental operation in parallel computing! 🚀

```
#include <stdio.h>
#include <stdlib.h>
#include <time.h>
#include <mpi.h>

int main() {
    MPI_Init(NULL, NULL);

    int rank;
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);

    // Each process gets different random numbers
    srand(time(NULL) + rank);
    int a = rand() % 100;
    int b = rand() % 100;
    int local_sum = a + b;

    printf("Process %d: %d + %d = %d\n", rank, a, b, local_sum);

    // Sum all local sums to process 0
    int global_sum;
    MPI_Reduce(&local_sum, &global_sum, 1, MPI_INT, MPI_SUM, 0, MPI_COMM_WORLD);

    if (rank == 0) {
        printf("\nTotal sum from all processes: %d\n", global_sum);
    }

    MPI_Finalize();
    return 0;
}
```

How to Use:

1. Save as `simple_mpi_sum.c`
2. Compile: `mpicc simple_mpi_sum.c -o simple_mpi_sum`
3. Run: `mpirun -n 4 ./simple_mpi_sum`

Sample Output:

```
Process 1: 28 + 65 = 93
```

```
Process 2: 14 + 70 = 84
```

```
Process 3: 53 + 41 = 94
```

```
Process 0: 67 + 32 = 99
```

```
Total sum from all processes: 370
```

Key Features:

- Each process generates just 2 random numbers (0-99)
- Calculates local sum
- Uses MPI_Reduce to sum all local sums
- Process 0 displays final result
- Clean and minimal - focuses just on the reduction operation