

Week13-mpi-with-python

Basic MPI Concepts in Python

1. **Ranks**

Each process in MPI gets a unique ID called a "rank".

`rank == 0` is usually the **master** process.

2. **Size**

Total number of processes running.

3. **Communication**

Use `send()`, `recv()`, or collective operations like `bcast()` and `gather()`.

Example: Hello World with `mpi4py`

```
# hello_mpi.py
from mpi4py import MPI

comm = MPI.COMM_WORLD
rank = comm.Get_rank()
size = comm.Get_size()

print(f"Hello from process {rank} of {size}")
```

Run it with 4 processes:

```
mpiexec -n 4 python hello_mpi.py
```

mpi4py

is a **Python wrapper** for the **MPI (Message Passing Interface)** standard, allowing Python programs to perform **parallel processing** across multiple CPUs, cores, or even multiple machines — just like traditional HPC (High Performance Computing) applications written in C or Fortran using MPI.

What is `mpi4py`?

- `mpi4py` is a Python package that **binds the MPI C API** to Python.
 - It enables you to:
 - Launch multiple processes (`mpiexec -n N python your_script.py`)
 - Send and receive messages (`send()`, `recv()`)
 - Use collective communication (`bcast()`, `scatter()`, `gather()`, `reduce()`)
 - It supports **NumPy arrays**, making it efficient for large data communication.
-

Core Concepts in `mpi4py`

Concept	MPI API	Description
World Comm	<code>COMM_WORLD</code>	The default communicator (all processes)
Rank	<code>Get_rank()</code>	Each process has a unique rank ID
Size	<code>Get_size()</code>	Total number of processes in the group
Send/Recv	<code>send()</code> , <code>recv()</code>	Point-to-point messaging
Broadcast	<code>bcast()</code>	Send data from one to all
Scatter	<code>scatter()</code>	Split data across processes
Gather	<code>gather()</code>	Collect data from all processes
Reduce	<code>reduce()</code>	Aggregate data (e.g., sum, max)
Barrier	<code>barrier()</code>	Synchronize all processes

Why Use `mpi4py`?

- Run the same code across multiple processes (not just threads).
- Avoid Python's Global Interpreter Lock (GIL) for CPU-bound tasks.
- Scale programs across multiple nodes in an HPC cluster.
- Works great with NumPy for numerical and matrix-heavy applications.

1. Point-to-Point Communication: `send()` and `recv()`

```
# point_to_point.py
from mpi4py import MPI

comm = MPI.COMM_WORLD
rank = comm.Get_rank()

if rank == 0:
    data = "Hello from rank 0"
    comm.send(data, dest=1)
    print(f"[{rank}] Sent data to rank 1")
elif rank == 1:
    data = comm.recv(source=0)
    print(f"[{rank}] Received data: {data}")
```

2. Broadcast: `bcast()`

```
# broadcast.py
from mpi4py import MPI

comm = MPI.COMM_WORLD
rank = comm.Get_rank()

data = None
if rank == 0:
    data = "Broadcast message"

data = comm.bcast(data, root=0)
print(f"[{rank}] Got: {data}")
```

3. Scatter: `scatter()`

```
# scatter.py
from mpi4py import MPI

comm = MPI.COMM_WORLD
rank = comm.Get_rank()
size = comm.Get_size()

if rank == 0:
    data = [i for i in range(size)] # One value per process
else:
    data = None

recv_data = comm.scatter(data, root=0)
print(f"[{rank}] Received {recv_data}")
```

4. Gather: `gather()`

```
# gather.py
from mpi4py import MPI

comm = MPI.COMM_WORLD
rank = comm.Get_rank()

send_data = rank * 2
gathered_data = comm.gather(send_data, root=0)

if rank == 0:
    print(f"[{rank}] Gathered data: {gathered_data}")
```

5. Reduce: `reduce()`

```
# reduce.py
from mpi4py import MPI

comm = MPI.COMM_WORLD
rank = comm.Get_rank()

send_value = rank + 1 # 1 for rank 0, 2 for rank 1, etc.
total = comm.reduce(send_value, op=MPI.SUM, root=0)

if rank == 0:
    print(f"[{rank}] Sum of ranks+1 = {total}")
```



Example: Distributed Array Sum

```
# sum_mpi.py
from mpi4py import MPI
import numpy as np

comm = MPI.COMM_WORLD
rank = comm.Get_rank()
size = comm.Get_size()

# Split array among processes
if rank == 0:
    data = np.arange(100, dtype='i')
    chunks = np.array_split(data, size)
else:
    chunks = None

# Scatter chunks to all processes
local_data = comm.scatter(chunks, root=0)
```

```
# Each process computes local sum
local_sum = np.sum(local_data)

# Gather all local sums at root
total_sum = comm.reduce(local_sum, op=MPI.SUM, root=0)

if rank == 0:
    print("Total sum:", total_sum)
```

When to Use MPI in Python?

- **Large-scale data processing**
 - **Scientific simulations**
 - **Parallel numerical computations**
 - Where **shared memory models (like multithreading)** won't scale across machines
-

Things to Keep in Mind

- MPI excels on clusters; not always efficient on a single machine.
- Python has overhead—use compiled extensions (NumPy, Numba) for heavy lifting.
- Debugging is harder—consider writing unit tests for components before running them distributed.

How to setup in windows?

Microsoft MPI v10.1.3

Stand-alone, redistributable and SDK installers for Microsoft MPI

<https://www.microsoft.com/en-us/download/details.aspx?id=105289>

 **To fix this on Windows, follow these steps:**

1. Install Microsoft MPI

Download the **Microsoft MPI (MS-MPI)** from the official site:

<https://www.microsoft.com/en-us/download/details.aspx?id=105289>

- Install both:
 - **MS-MPI Redistributable Package**
 - **MS-MPI SDK**

This will install `mpiexec.exe` and necessary DLLs.

2. Add MPI to System PATH (if needed)

If after install it still fails, manually add this to your PATH:

`C:\Program Files\Microsoft MPI\Bin`

3. Verify Installation

Open a new **Command Prompt**, and run:

`mpiexec`

You should see usage info instead of an error.

4. Install `mpi4py` (if not done yet)

```
pip install mpi4py
```

 **Test It**